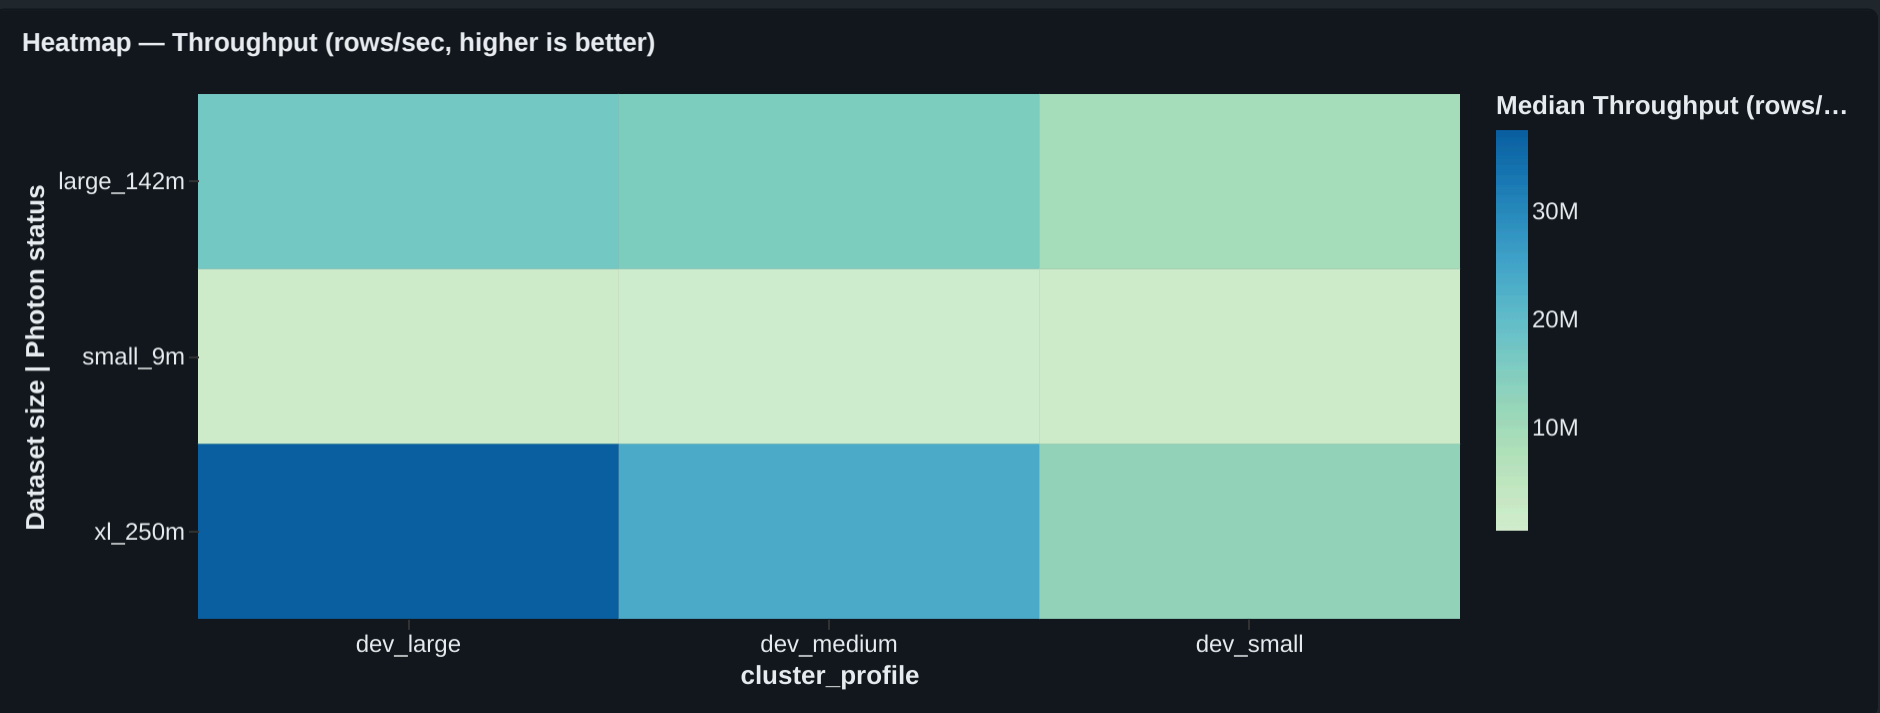

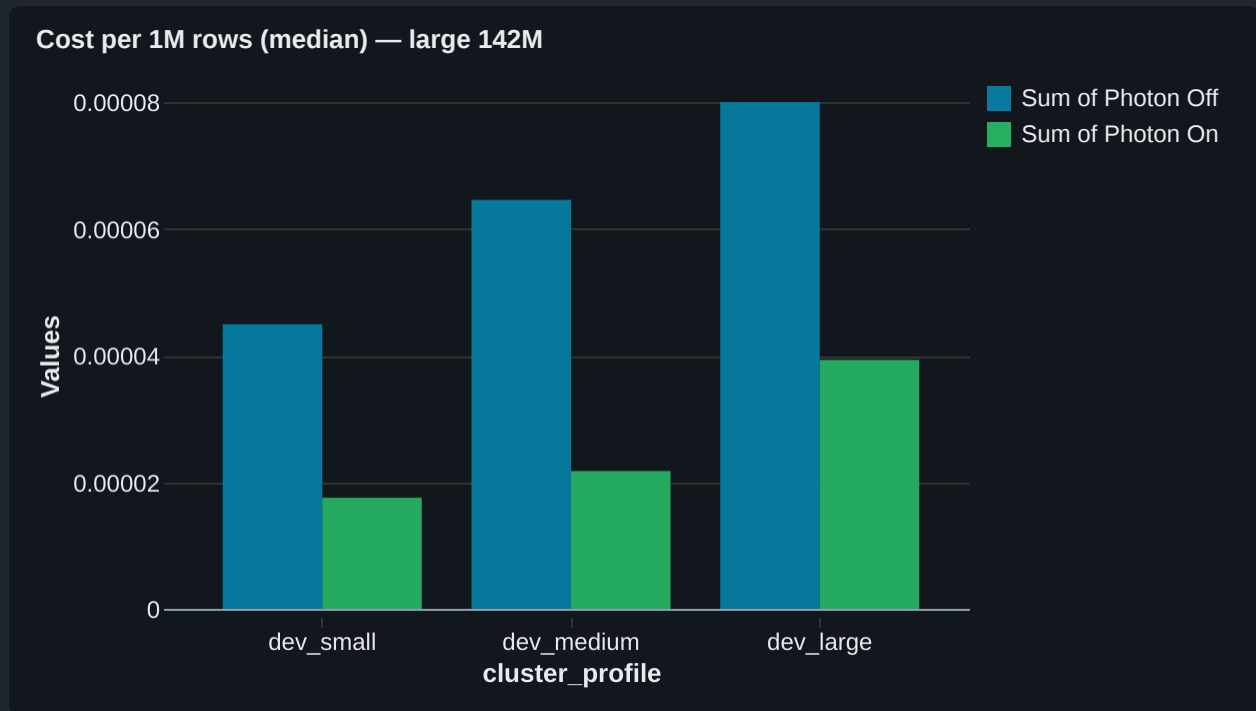
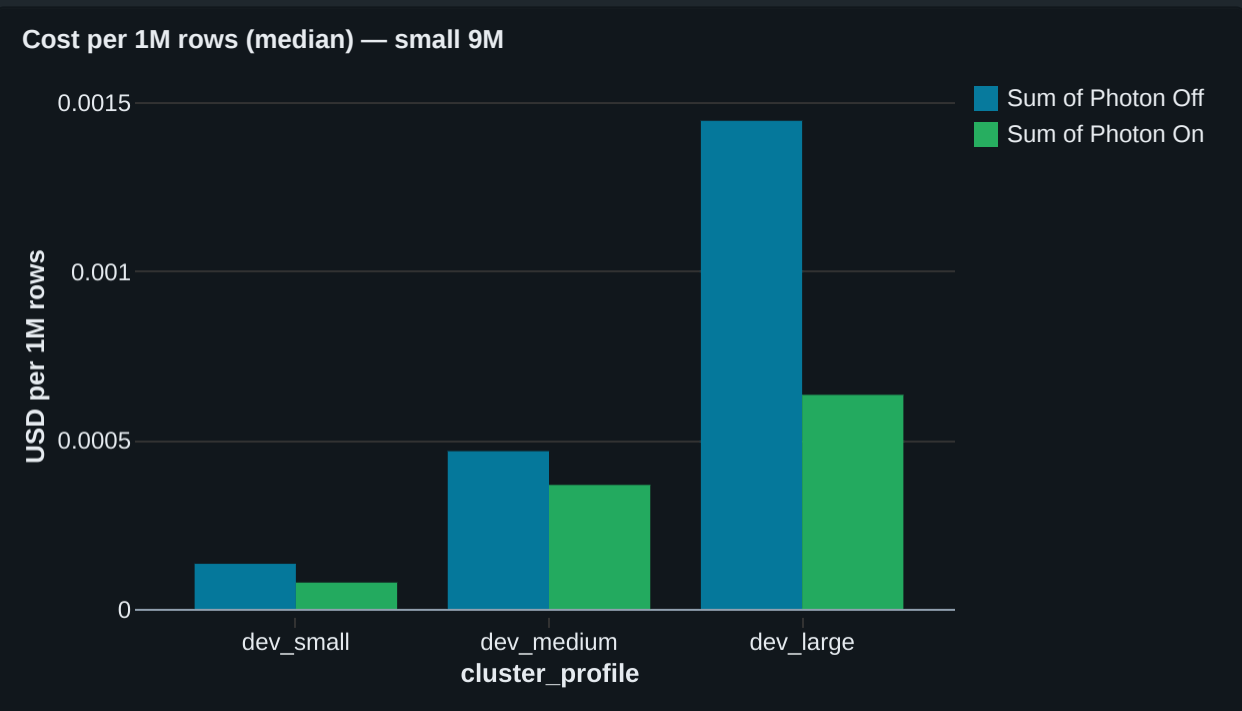
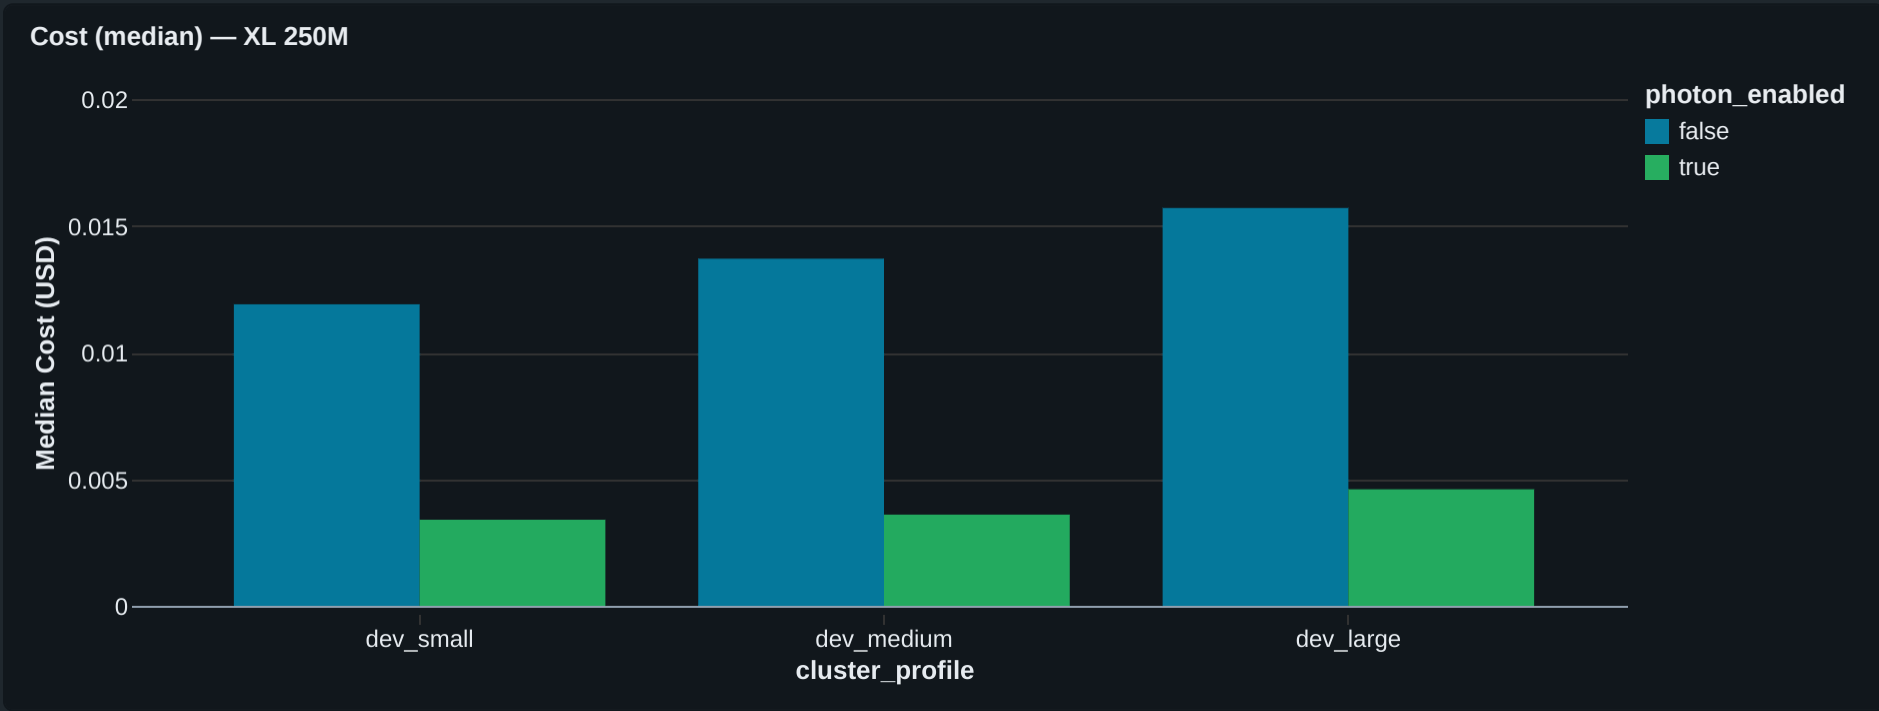
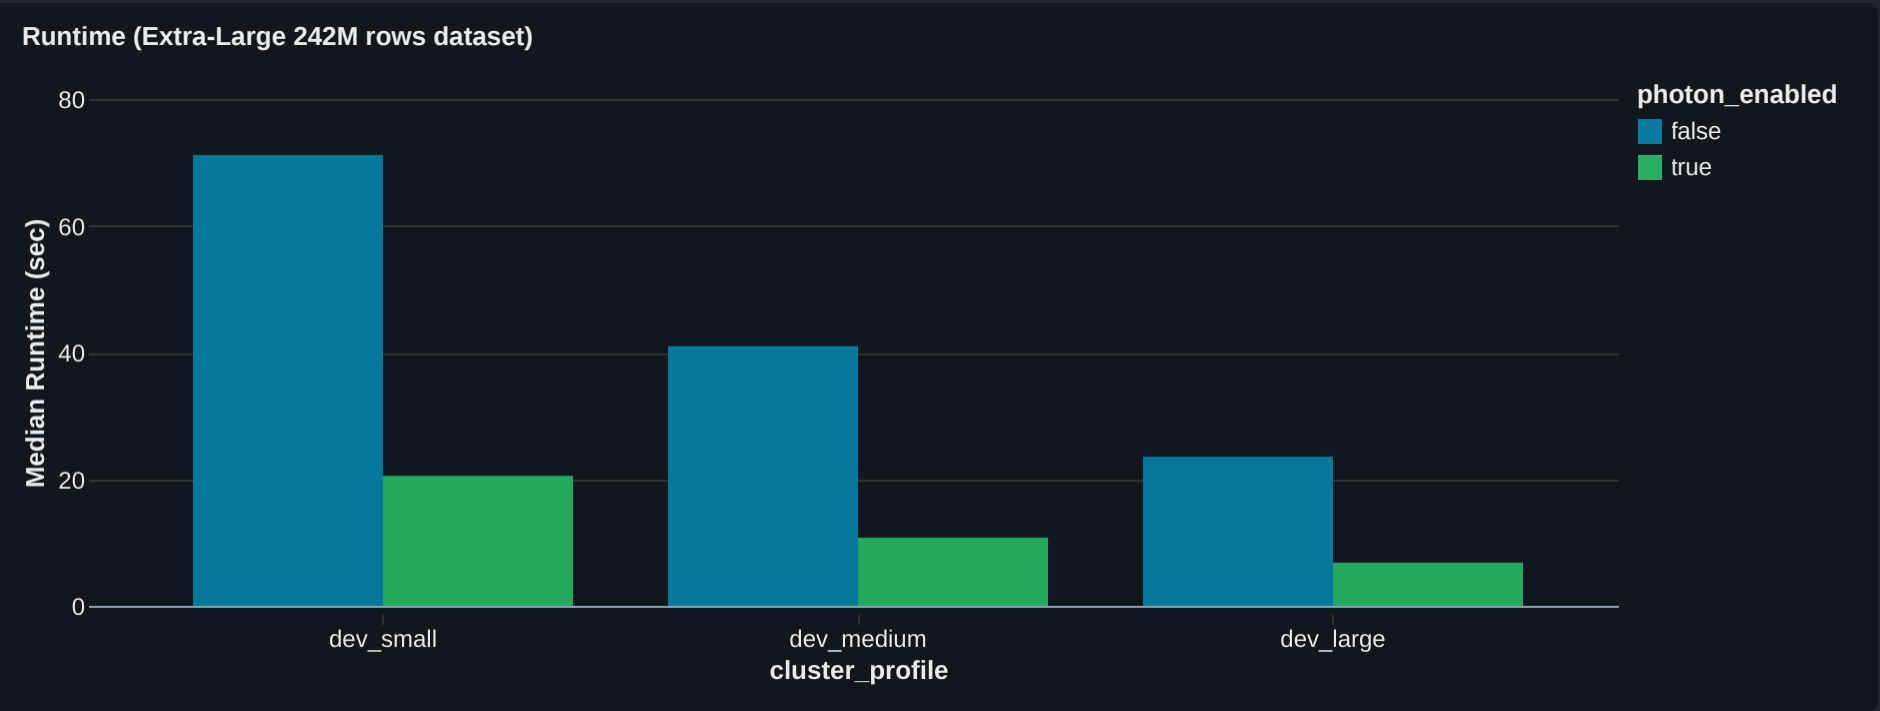

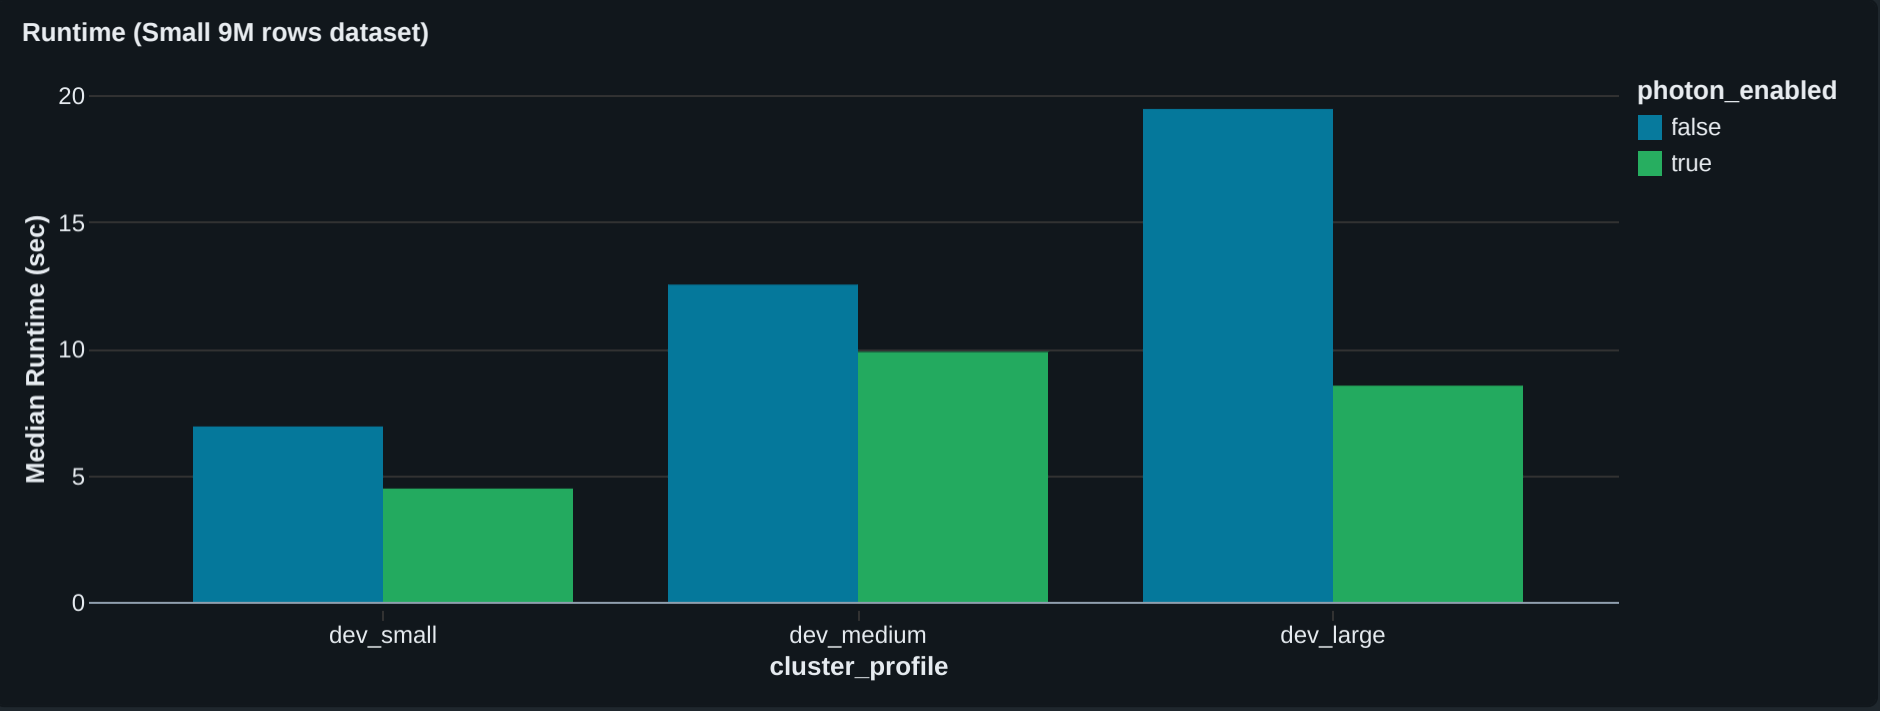


Delta Pipeline Benchmarking Results (9M -> 250M rows)



Key Takeaways

- Photon reduces median runtime by ~2–4× and cost by ~3–4× on large and XL datasets.
- Best value choice: dev_medium + Photon → balances speed and cost across scales.
- When speed is priority: dev_large + Photon delivers fastest runtimes (especially XL).
- For small datasets: dev_small remains cheapest; Photon still helps, but gains are modest.
- Cost per 1M rows drops significantly with Photon, confirming scale-driven efficiency.
- (Note: dev_small = 2 workers, dev_medium = 4 workers, dev_large = 8 workers)